

CREDIT CARD FRAUD DETECTION

- PINAKI BHAGAT
- SYDNEY CORREA

Learning Unbalanced Data

Agenda

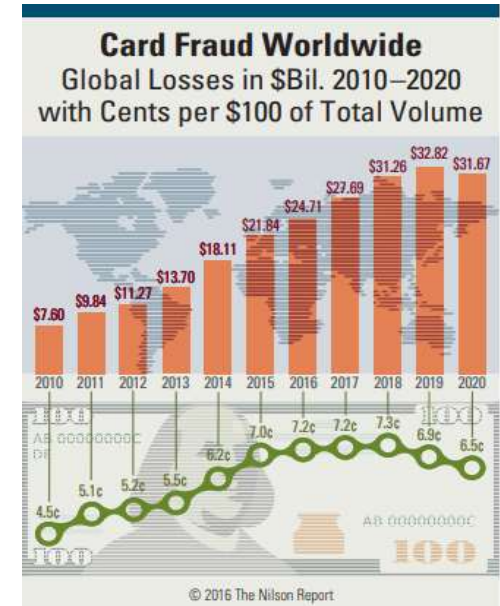


- Credit Card use case overview
- Dataset description
- Techniques: Traditional/Deep Learning/ML Platforms
- Findings
- Future enhancements

Overview

Credit card fraud detection (Learning unbalanced data)

- Annual global fraud losses reached \$21.8 billion in 2015.
- Approx 12 cents per \$100 were stolen in the US during the same year
- Future trends show moving toward fraud prediction (and prevention) rather than fraud detection
- The cost of false positive fraud labelling was \$118 billion dollars of which \$9 billion were actual fraud cases
- Trends are moving from more traditional approaches to deep learning approaches (including hybrids w/ graph dbs)



Project objective

Objective: Try different approaches to compare model performance

A. Traditional methods

- a. K-means, Random Forest

- b. Ensemble Learning: Random forest, ExtraTrees, AdaBoost, GradientBoost, XGBoost

B. Deep Learning

- a. Autoencoders

C. ML Platforms

- a. h2o.ai


Research some of the leading approaches and cutting edge technologies

Credit card Dataset

The creditcard.csv datasets is one of 4 datasets on Kaggle and contains transactions made by credit cards in September 2013 by European cardholders.

Features:

- transactions that occurred in two days
- contains 492 frauds out of 284,807 transactions
- highly unbalanced (frauds account for 0.172% of all transactions)
- dataset contains only numerical input variables which are the result of a PCA transformation.
- features “Time” and “Amount” are not transformed



Techniques: KNN, RF

Classification problem using shallow algorithms to train individual models using KNN and RF

Checked for correlation between features

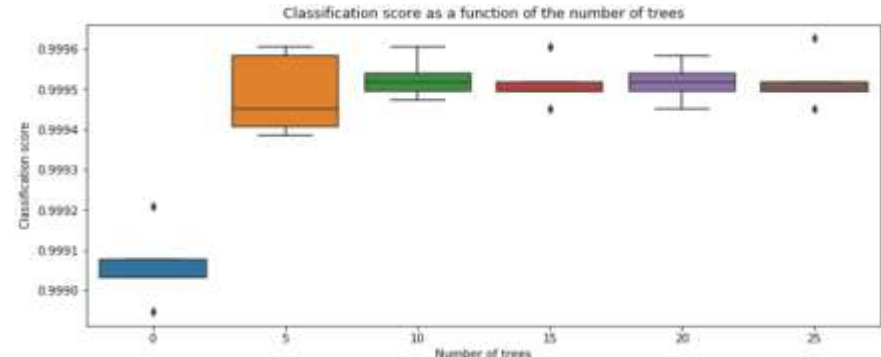

Hyper tuned the RF model for $n_estimators$ ($n=16$)

Ran KNN for different intervals ($k=5$)

Checked model performance using:

- AUPRC (area under precision recall curve)
- ROC
- Confusion matrix

The optimal number of neighbors is 5



Techniques: Ensemble

Following algorithms were used to train the model stack:

- Random Forest
- Extra Trees
- AdaBoost
- GradientBoost

Co-relation between models were checked and extra trees was eliminated

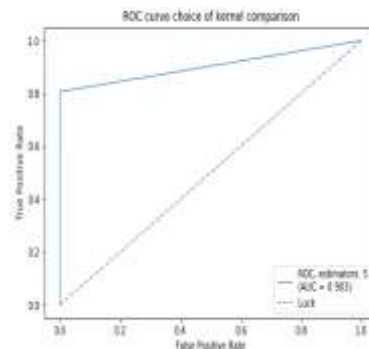
Ensemble had XGBoost as the final model

Accuracy score: 99.95%

Classification report:

	precision	recall	f1-score	support
Output 0	1.00	1.00	1.00	113726
Output 1	0.93	0.76	0.83	197
avg / total	1.00	1.00	1.00	113923

	AdaBoost	ExtraTrees	GradientBoost	RandomForest
0	0.278610	0.000171	0.000346	0.000128
1	0.042268	0.008257	0.000353	0.000136
2	0.280778	0.000221	0.000360	0.000144
3	0.291032	0.000237	0.000376	0.000304
4	0.276816	0.000298	0.000363	0.000098




Techniques: Autoencoders

Anomaly detection by training the network on non-fraudulent transactions, fraudulent transactions to be detected by flagging them based on higher than normal MSE values (reconstruction error)


Features:

- 4 fully connected layers (two encoders and 2 decoders)
- Activation filters: tanh, relu


Reconstruction error for Autoencoder: $L(x, x') = ||x - x'||^2$



Reconstruction error: Class 0



Reconstruction error: Mixed



Techniques: h2o.ai

We explored using the popular H2O.ai platform, specifically the AutoML package.

AutoML packages – Random Forest (DRF, XRT), XGBoost GBM, Deep Neural Net, ...

Hyperparameter tuning – max_models (5,10,15), balancing

```
aml = H2OAutoML(max_models=10, seed=1234, balance_classes = True, max_after_balance_size=0.8)  
aml.train(x = predictors, y = response, training_frame = train, validation_frame = valid)
```

	model_id	auc	logloss	mean_per_class_error	rmse	mse
	XGBoost_1_AutoML_20181219_234301	0.984103	0.00250716	0.101853	0.0200578	0.000402316
	GLM_grid_1_AutoML_20181219_234301_model_1	0.973093	0.00411683	0.100894	0.0265999	0.000707553
	XGBoost_2_AutoML_20181219_234301	0.971338	0.00324639	0.10298	0.0203073	0.000412386
	DRF_1_AutoML_20181219_234301	0.956017	0.00393489	0.0894407	0.0207101	0.000428908
	XRT_1_AutoML_20181219_234301	0.952861	0.00359792	0.10298	0.0206141	0.000424941
	StackedEnsemble_BestOfFamily_AutoML_20181219_234301	0.94077	0.00311164	0.0950754	0.0207072	0.000428788
	StackedEnsemble_AllModels_AutoML_20181219_234301	0.93763	0.00312438	0.09733	0.020708	0.000428822



Sources: h2o.ai

Techniques: Balancing

Considering the nature of dataset transactions to be unbalanced (0.172%) we tried multiple under sampling and oversampling techniques

Packages tried:

Oversampling - imblearn (SMOTE, SMOTEEEN);

Undersampling – RandomUnderSampler, H2o.ai


Models tried: KNN and RF and h2o AutoML

Does not apply to Autoencoders and

Outcomes: Model performance was significantly worse with balanced data.



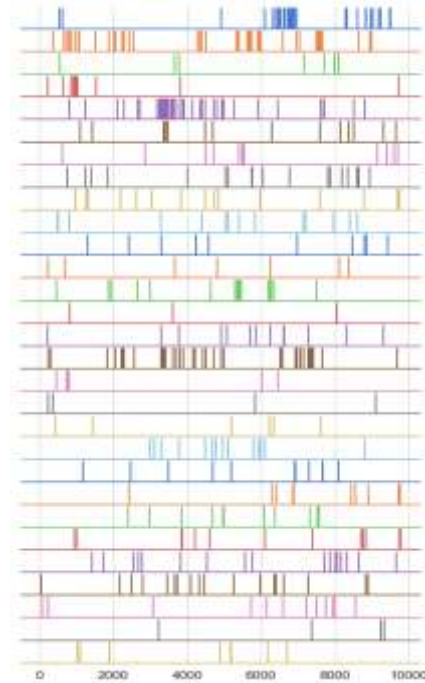
The optimal number of neighbors is 1



Pic Source: freepik.com

Findings

Fraud transactions: ridge line graph using seaborn



Future enhancements



Citation & Useful links

- ❑ <https://www.slideshare.net/databricks/deep-learning-for-recommender-systems-with-nick-pentreath>
- ❑ https://github.com/maciejkula/spotlight/tree/master/examples/movielens_explicit
- ❑ <https://recsys.acm.org/recsys18/tutorials/#content-tab-1-0-tab>
- ❑ <https://code.fb.com/core-data/recommending-items-to-more-than-a-billion-people/>
- ❑ <https://www.fast.ai/2017/07/28/deep-learning-part-two-launch/>
- ❑ <https://www.ethanrosenthal.com/2015/11/02/intro-to-collaborative-filtering/>
- ❑ <https://towardsdatascience.com/various-implementations-of-collaborative-filtering-100385c6dfe0>
- ❑ <https://www.youtube.com/watch?v=h9gpufJFF-0>
- ❑ https://en.wikipedia.org/wiki/Collaborative_filtering